

CIS 3207 Project 4

Designing and Implementing a File System as a Virtual Disk

1 Introduction

Operating systems need to provide a uniform, logical view of information storage. The operating system abstracts from the physical properties of its storage devices to define a logical storage unit, **the file**. Files are mapped by the operating system onto physical devices. A file is simply a *named collection of related information* that is recorded on secondary storage.

The information in a file is defined by its creator. A file may have a certain defined structure, which depends on its type, or the type of data stored in the file. The data may also be unstructured, i.e., simply a sequence of bytes.

A **file system** is comprised of a sequence of bytes allocated as a storage container on a physical storage device. The file system consists of two distinct parts: a collection of files, and a directory structure that organizes and provides information about all of the files in the system. The storage container is comprised of data within the files, and meta data that describes the files (the file system directory). Through the file system directory, a user has access to files, the file descriptions, the location of file content and the file content itself.

This project has been developed to help you understand the concepts of a file system, the relationship between files and the file system, and the requirements and structure of a directory of a file system. In this project, you will design a file system and implement it on a disk within the Linux operating system. A good introduction to file systems is in your Three Easy Pieces textbook, Chapters 39 and 40. Information about mass storage/disk systems and I/O management can be found in Chapters 36 and 37. Additional references are given at the end of this document.

Generally, a file system encompasses all of the storage of a device. That is, we generally consider a storage device, such as a disk, as a contiguous set of bytes, and all of the bytes are available to the file system. The file system gives structure to the storage elements (bytes) on the device. In fact, there is a *logical view* for the contents of a file, and a *physical view*. One can consider that the file system provides the mapping between these views. In this project, “**your file system**” will be **created within a single file** (allocated storage of bytes) of the operating system disk on which you are developing the file system. That is, you are creating a ‘virtual disk’, implemented within a file stored on the operating system disk. (Remember, a file is just a sequence of bytes; you can impose any data structure(s) on this sequence of bytes. In this project, the sequence of bytes of this file will comprise your entire file system).

For the project, you will allocate a large file on your system disk (one large enough to hold both the directory and all the files to be stored in this file system) that can be **logically viewed** as a storage device composed of contiguous bytes. We can call this file a ‘virtual disk’. You will give structure and organization to those bytes. You will create a directory of files that must include a component that describes the allocation of bytes to files and includes information about each file (name, creation date, size, etc.). In the filesystem, bytes are allocated to files and to the directory system itself; i.e., all bytes that are within the disk must be accounted for and tracked. The directory will be a “file” or allocated storage area within your storage container (i.e., in your file system). Each entry in the directory will be a data structure describing the file it represents.

The file that you use as your virtual disk does not have to be physically contiguous (i.e., bytes adjacent to one another). The underlying operating system (e.g., Linux) takes care of the physical allocation of all files, making them appear to be logically contiguous. You can create a large file on the operating system's disk and ignore the real physical organization of the bytes in the file, viewing the storage as **logically contiguous** (your view of the disk you are creating). You can create the virtual disk file and write bytes for the length of the file you will create.

That is, a starting point for this project is to create a large file on a disk. This file represents "your disk", or virtual device. Your objectives in this project are to create the file system for that virtual device and develop the basic file operations for the device. These operations will include creating and deleting files, opening and closing files, and viewing the directories for your file system.

2 Your Tasks

1. Create a Filesystem
 - a. This file system is to have a directory structure that can store (sub)directories and files.
 - b. Files and directories can have names with as few as 8 characters along with a 3-character filename extension (file type descriptor). (this is a minimum requirement; you could have longer filenames)
 - c. Files can be up to 65536 characters (bytes) in size [again, a minimum requirement.]
 - d. Storage of individual files should be designed such that non-contiguous files are stored in the filesystem.
 - e. The allocation units, or blocks, on the disk you are to create are each 512 bytes.
 - f. Directory structure
 - i. The directory is to be extensible; i.e., all the elements in the directory can be linked (this does not imply by a linked list) so that additional file/directory entries can be added, or entries removed when files are deleted.
 - ii. Every entry in a directory is to have a file name, creation date and time, last modification date/time stamp, file size and information about the space allocated to the file or a link to such information,.
 - g. You are to select and design the specific details of the layout of a directory, including the root directory. Your design should include:
 - i. Structure and organization of the directory information
 - ii. Structure of the file identification in the directory
 - iii. Structure of the location of file components (i.e., physical location of file blocks)
 - iv. Method of describing the full path to a file.
 - h. The file system is to have a **logical directory** that describes the files and meta data as noted above, and a **physical directory** that organizes the blocks or disk units allocated to each file. The specifics of these directories and allocation methods are to be of **your design**.
 - i. The physical directory will keep track of the blocks of disk space allocated to your disk itself, as well as each file in the system.
2. Create functions to control file I/O within the filesystem
 - a. Create File
 - b. Delete File
 - c. Open File
 - d. Close File

3. All I/O in the file system should be performed using “memory=mapped file” functionality, where your filesystem is mapped into memory and manipulation of the filesystem (or virtual disk) is directly through memory
4. You are to design testing procedures to demonstrate the functional and physical organization of the directory structure as it expands and contracts with the addition and deletion of files.
 - a. This means that you need to be able to create files of various sizes and add them to your disk and remove them (deletion).
 - b. Note that to accomplish the above, you will need to be able to read or write to space in your filesystem and well as the filesystem’s directory itself.
5. With **memory mapped files**, you open the virtual disk file and map the file to memory using the `mmap()` system call and facility. The memory mapped file facility lets you address locations in the file and read and write the locations as if you are directly addressing memory locations. In effect, the various locations in your virtual disk are mapped to memory cells. Reading and writing to disk is managed by the OS as you read and write the memory cells. There are many references for `mmap` programming, including:

[Linux System Programming](#) (Love)

The Linux man pages for `mmap`

6. Some References for File Systems currently in Use

- a. File System Wiki
http://en.wikipedia.org/wiki/file_system#File_systems_under_Unix-like_operating_systems
- b. FAT File System:
 - i. Design of the FAT file system
https://en.wikipedia.org/wiki/Design_of_the_FAT_file_system
 - ii. **FAT Wiki** http://en.wikipedia.org/wiki/File_Allocation_Table
 - iii. <http://www.ntfs.com/fat-systems.htm>
 - iv. http://www.pctechguide.com/31HardDisk_File_systems.htm
 - v. <http://msdn.microsoft.com/en-us/library/aa914353.aspx>
- c. NTFS File System
 - i. <http://www.ntfs.com/>
 - ii. <http://en.wikipedia.org/wiki/NTFS>
- d. Unix File System
 - i. Unix Fast File System (<https://www.cs.berkeley.edu/~brewer/cs262/FFS.pdf>)
 - ii. Anatomy of Linux journaling file systems
(<http://www.ibm.com/developerworks/library/l-journaling-file-systems/>)
 - iii. Next-generation Linux file systems: NiLFS(2) and exofs
(<http://www.ibm.com/developerworks/linux/library/l-nilfs-exofs/index.html>)
- e. General Information on a File System Design
 - i. Practical File System Design
 1. (<http://www.nobius.org/~dbg/practical-file-system-design.pdf>)